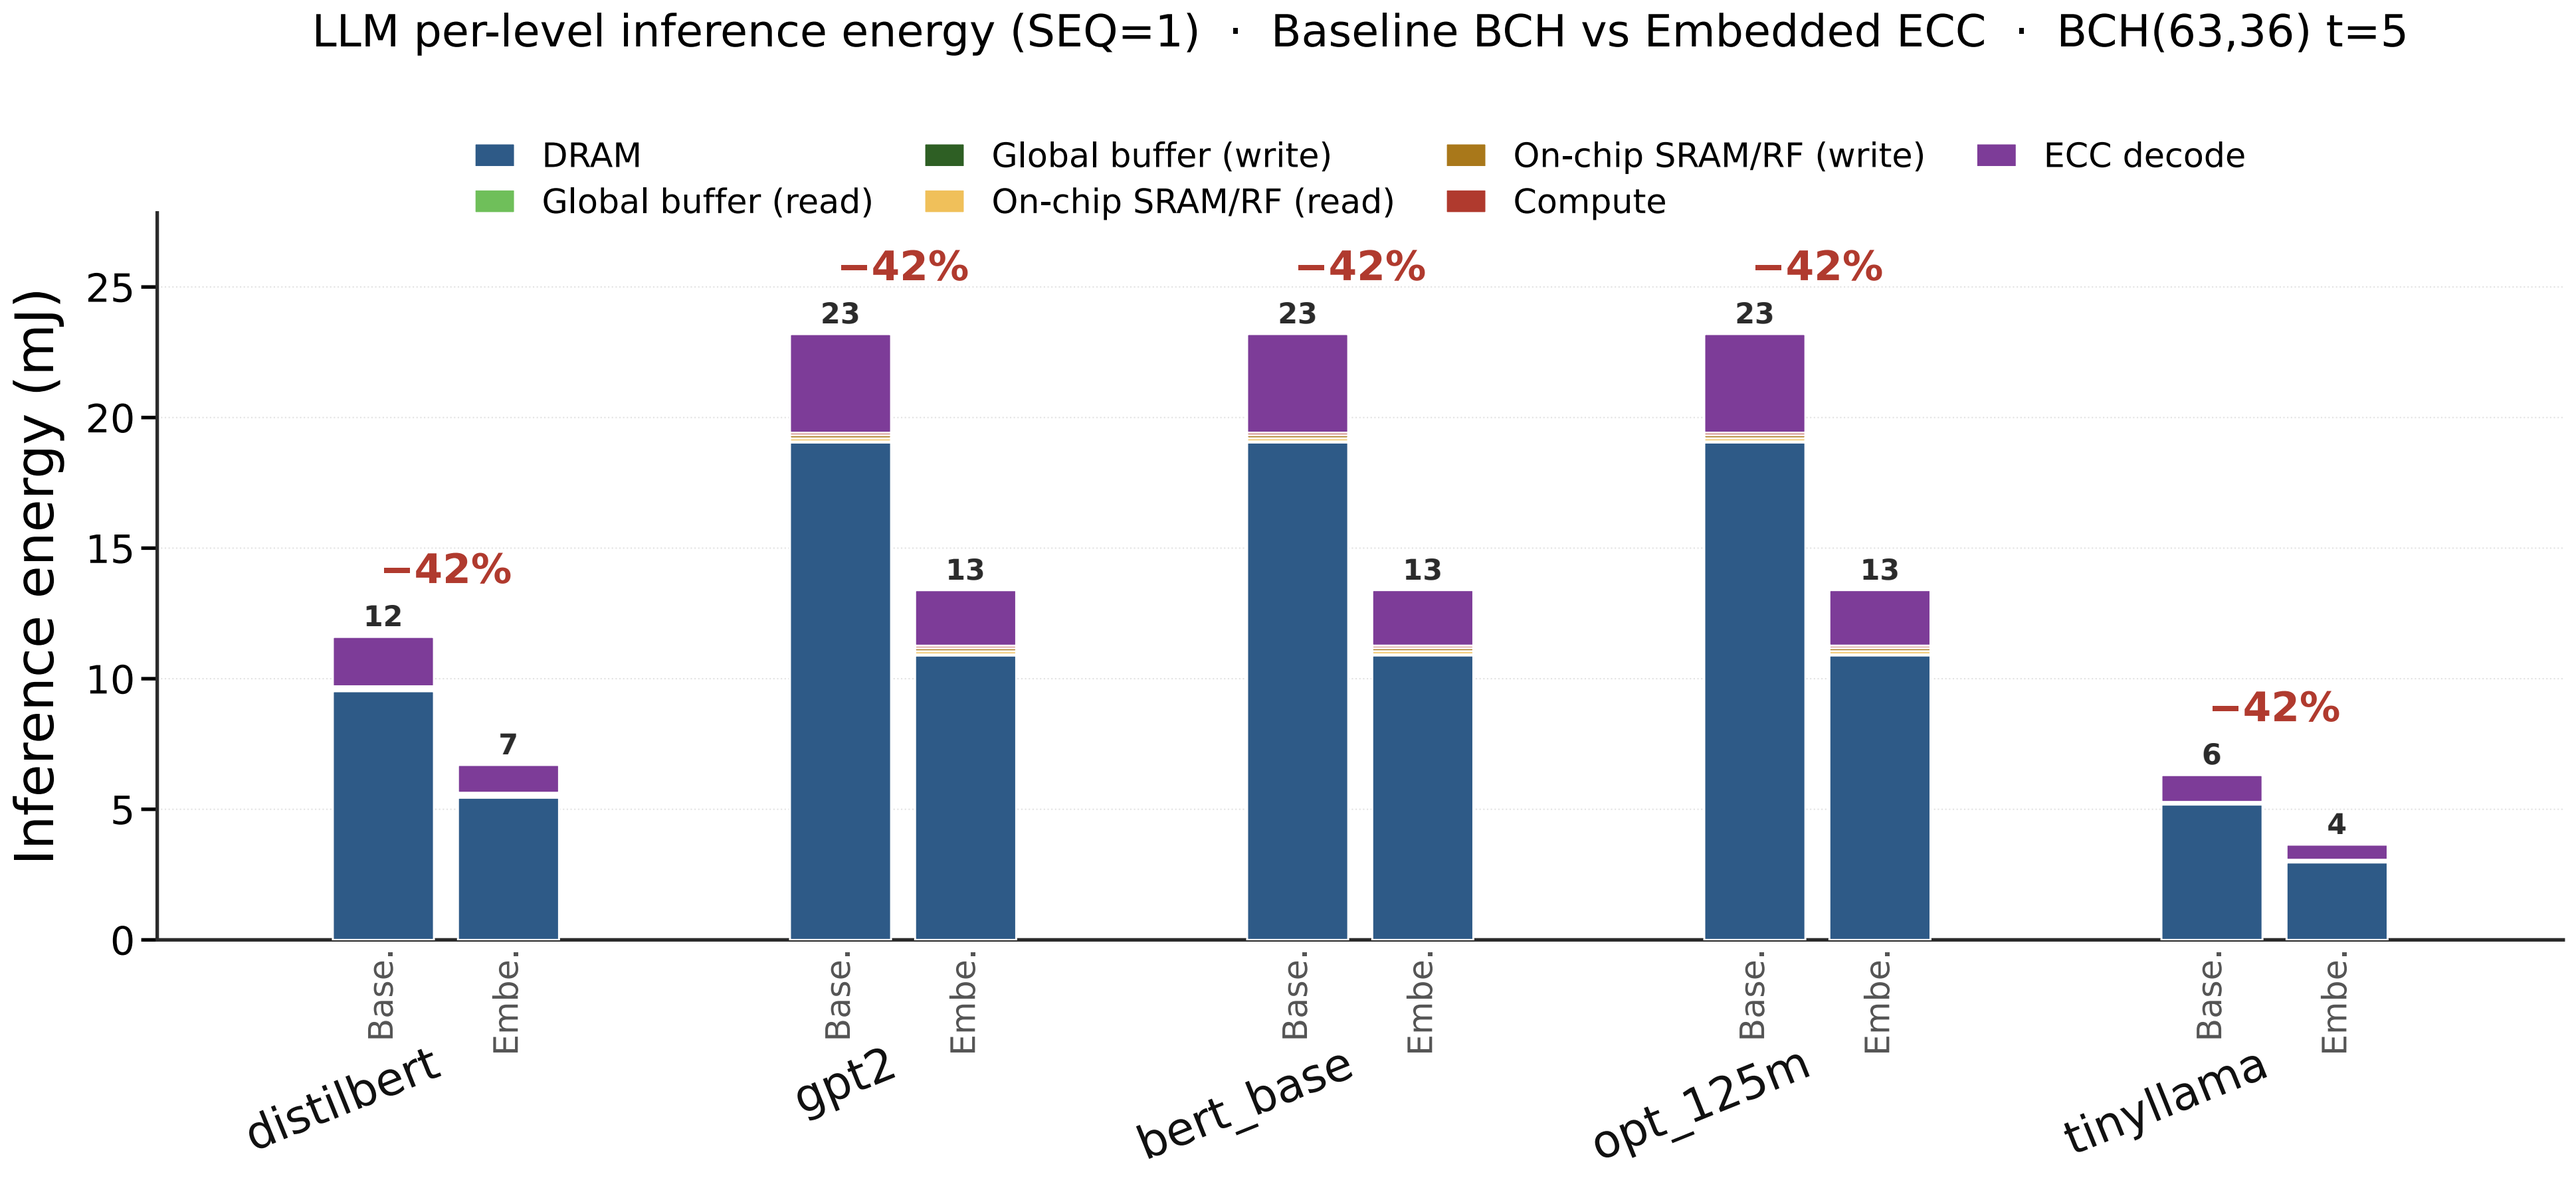

LLM per-level inference energy (SEQ=1) · Baseline BCH vs Embedded ECC · BCH(63,36) t=5
DRAM
Global buffer (read)
Global buffer (write)
On-chip SRAM/RF (read)
On-chip SRAM/RF (write)
Compute
ECC decode
Inference energy (mJ)
0
5
10
15
20
25
−42%
12
7
−42%
23
13
−42%
23
13
−42%
23
13
−42%
6
4
Base.
Embe.
distilbert
gpt2
bert_base
opt_125m
tinyllama